Praveen Madupu
Mb: +91 98661 30093
Sr. Database Administrator

Real-time scenario-based interview questions for **DB2 LUW (Linux, Unix, Windows) Administration, Performance Tuning & Optimization:**

🟢 DB2 LUW Administration, Performance Tuning & Optimization – 500 Real-Time Scenario Questions

🔷 Basic Level (0–2 years) – 1–100

Installation & Configuration

1. How do you install DB2 LUW on Linux and Windows for production environments?
2. How do you configure instance-level parameters during DB2 installation?
3. How do you configure DB2 memory parameters for optimal startup performance?
4. How do you change the default DB2 instance and database directories?
5. How do you enable and configure DB2 audit logging?
6. How do you configure DB2 to start automatically on system boot?
7. How do you enable TCP/IP connectivity for DB2 databases?
8. How do you configure DB2 communication ports for multiple instances?
9. How do you secure the DB2 configuration file (db2profile)?
10. How do you verify DB2 installation and check for patch levels?

Instance & Database Management

11. How do you create a new DB2 instance and verify it?
12. How do you delete a DB2 instance safely?
13. How do you create a new DB2 database with custom storage paths?
14. How do you drop a DB2 database safely in production?
15. How do you rename a DB2 database?
16. How do you attach a user to a DB2 instance?
17. How do you update DB2 database configuration parameters online?
18. How do you monitor instance-level configuration parameters for performance?
19. How do you list all active databases on an instance?
20. How do you enable logging for database creation and deletion events?

User Management & Security

21. How do you create DB2 users with role-based privileges?
22. How do you grant read-only access to specific schemas?
23. How do you revoke privileges safely in production?
24. How do you enforce strong password policies for DB2 users?
25. How do you detect unauthorized login attempts in DB2?
26. How do you audit user activities using DB2 audit facility?
27. How do you integrate LDAP authentication with DB2?
28. How do you rotate DB2 instance passwords without downtime?
29. How do you implement connection restrictions using IP or hostname?
30. How do you configure DB2 to enforce SSL/TLS connections?

Praveen Madupu
Mb: +91 98661 30093
Sr. Database Administrator

Backup & Recovery Basics

31. How do you take a full backup of a DB2 database?
32. How do you restore a DB2 database from a full backup?
33. How do you perform incremental or delta backups?
34. How do you automate backup scheduling for production databases?
35. How do you backup only selected tablespaces?
36. How do you verify the integrity of a DB2 backup?
37. How do you compress backups to save storage?
38. How do you restore a single tablespace from backup?
39. How do you perform point-in-time recovery using logs?
40. How do you implement backup strategies for high availability environments?

Monitoring & Logging Basics

41. How do you enable and analyze DB2 diagnostic logs?
42. How do you monitor buffer pool hit ratios?
43. How do you monitor active connections and sessions?
44. How do you detect high CPU-consuming SQL statements?
45. How do you track long-running transactions?
46. How do you detect and troubleshoot deadlocks?
47. How do you monitor index usage in production?
48. How do you track disk I/O per tablespace or database?
49. How do you monitor lock contention?
50. How do you enable event monitors for performance tracking?

Tables & Index Management

51. How do you create a table with primary and foreign keys?
52. How do you alter a table to add a new column without downtime?
53. How do you rename a table safely in production?
54. How do you detect unused or redundant indexes?
55. How do you implement composite indexes for multi-column queries?
56. How do you rebuild fragmented indexes safely?
57. How do you monitor tablespace growth over time?
58. How do you detect table fragmentation and reclaim space?
59. How do you implement table partitioning for large tables?
60. How do you migrate a table from one tablespace to another?

Query Performance Basics

61. How do you detect slow queries using db2exfmt or db2pd?
62. How do you use EXPLAIN to analyze SQL execution plans?
63. How do you optimize queries performing table scans?
64. How do you detect queries causing high buffer pool usage?
65. How do you create and use indexes to speed up queries?
66. How do you monitor frequently executed queries?

67. How do you detect queries returning excessive rows?
68. How do you optimize JOIN operations on large tables?
69. How do you optimize queries on columns with large VARCHAR/CLOB data?
70. How do you rewrite queries to avoid correlated subquery issues?

Storage & Memory Basics

71. How do you monitor buffer pool usage?
72. How do you tune buffer pools for read-heavy workloads?
73. How do you monitor tablespace I/O?
74. How do you configure tablespaces with appropriate page sizes?
75. How do you detect memory pressure in DB2 instance?
76. How do you optimize temporary table usage?
77. How do you monitor and tune the sort heap?
78. How do you detect and prevent table or index fragmentation?
79. How do you reclaim unused tablespace?
80. How do you configure multiple data paths for large storage environments?

Security Basics

81. How do you implement SSL connections for client-server communications?
82. How do you audit user login attempts?
83. How do you revoke unnecessary privileges?
84. How do you enable encryption for data at rest?
85. How do you detect brute-force login attempts?
86. How do you implement firewall restrictions for DB2 connections?
87. How do you monitor privilege escalation attempts?
88. How do you rotate SSL/TLS certificates for the instance?
89. How do you secure DB2 configuration files?
90. How do you implement two-factor authentication for DB2 users?

Basic Replication & HA

91. How do you configure HADR (High Availability Disaster Recovery)?
92. How do you monitor HADR synchronization status?
93. How do you switch HADR roles (primary/standby)?
94. How do you detect replication failures or network issues?
95. How do you perform a manual takeover in HADR?
96. How do you add a new standby database to an existing HADR setup?
97. How do you remove a failed standby database safely?
98. How do you implement read-only standby for reporting?
99. How do you monitor log shipping status?
100. How do you ensure consistency between primary and standby databases?

Praveen Madupu
Mb: +91 98661 30093
Sr. Database Administrator

🔷 Intermediate Level (3–6 years) – 101–200

Query Performance & Optimization

101. How do you optimize queries with multiple joins?
102. How do you rewrite subqueries to joins for better performance?
103. How do you analyze EXPLAIN plans for bottlenecks?
104. How do you optimize queries with NULL values?
105. How do you detect queries causing temporary table spills?
106. How do you implement covering indexes for queries?
107. How do you detect queries causing sorts on disk?
108. How do you reduce disk I/O for large SELECT queries?
109. How do you optimize GROUP BY queries?
110. How do you handle large result sets efficiently?
111. How do you tune ORDER BY queries?
112. How do you optimize LIMIT/OFFSET queries?
113. How do you implement query hints in DB2?
114. How do you detect and resolve deadlocks?
115. How do you optimize queries on CLOB/BLOB columns?
116. How do you implement partition pruning for queries?
117. How do you monitor long-running transactions?
118. How do you analyze query cache hit ratios?
119. How do you optimize queries with IN or OR conditions?
120. How do you design queries to minimize lock contention?

Indexing Intermediate

121. How do you create multi-column indexes?
122. How do you detect unused indexes?
123. How do you rebuild fragmented indexes without downtime?
124. How do you optimize indexes for read-heavy workloads?
125. How do you optimize indexes for write-heavy workloads?
126. How do you detect redundant indexes?
127. How do you monitor index usage?
128. How do you implement partial or filtered indexes?
129. How do you analyze index selectivity?
130. How do you optimize indexes for foreign key lookups?
131. How do you implement indexes for partitioned tables?
132. How do you balance indexing for reads vs writes?
133. How do you use index-only access plans?
134. How do you detect queries ignoring indexes?
135. How do you implement indexes for frequently grouped queries?
136. How do you monitor index rebuild impact?
137. How do you detect contention on hot indexes?
138. How do you implement compressed indexes?

139. How do you optimize multi-table joins using indexes?
140. How do you detect and remove duplicate indexes?

Backup & Recovery Intermediate

141. How do you implement incremental backups efficiently?
142. How do you monitor backup job performance?
143. How do you restore individual tablespaces?
144. How do you perform point-in-time recovery using logs?
145. How do you automate backup verification?
146. How do you implement backup rotation policies?
147. How do you backup large databases efficiently?
148. How do you monitor backup storage usage?
149. How do you integrate backups with cloud storage?
150. How do you ensure consistent backups in HADR setups
151. How do you restore a corrupted backup?
152. How do you perform offsite backup automation?
153. How do you monitor backup job failures?
154. How do you optimize backup for minimal I/O impact?
155. How do you compress and encrypt backups?
156. How do you implement hot backups with minimal downtime?
157. How do you restore databases to alternate locations?
158. How do you handle large transaction logs during restore?
159. How do you implement backup strategies for multi-instance environments?
160. How do you monitor backup job duration trends?

Intermediate Administration, Monitoring, HADR

161. How do you tune DB2 memory parameters for high-concurrency workloads?
162. How do you monitor CPU and memory usage at instance-level?
163. How do you detect and resolve lock contention?
164. How do you monitor tablespace I/O statistics?
165. How do you monitor buffer pool hit ratios for performance tuning?
166. How do you configure HADR synchronous vs asynchronous modes?
167. How do you detect and troubleshoot HADR role switching failures?
168. How do you automate failover testing in HADR setups?
169. How do you implement read-only standby databases?
170. How do you optimize log shipping for minimal network impact?
171. How do you monitor active connections and transaction rates?
172. How do you implement connection throttling?
173. How do you detect resource-intensive queries?
174. How do you implement automated alerting for replication lag?
175. How do you tune sort heap size for complex queries?
176. How do you monitor lock waits and deadlock chains?
177. How do you detect table or index fragmentation in production?

Praveen Madupu
Mb: +91 98661 30093
Sr. Database Administrator

178. How do you optimize temporary tables for high transaction loads?
179. How do you monitor long-running utility operations?
180. How do you tune DB2 parameters for batch vs OLTP workloads?
181. How do you detect database-level I/O hotspots?
182. How do you implement automated scripts for performance stats collection?
183. How do you detect tables or indexes with skewed data distribution?
184. How do you tune buffer pools for partitioned tables?
185. How do you monitor replication consistency across multiple instances?
186. How do you configure alerts for high table or index growth?
187. How do you detect queries causing excessive memory consumption?
188. How do you monitor temp tablespace usage per session?
189. How do you detect long-running transactions blocking other queries?
190. How do you monitor deadlock chains and blocked transactions?
191. How do you tune instance-level memory pools for mixed workloads?
192. How do you detect performance degradation after DB2 upgrades?
193. How do you implement query monitoring for SLA adherence?
194. How do you detect sudden spikes in I/O or CPU usage?
195. How do you optimize stored procedures for large datasets?
196. How do you monitor package cache usage for repeated SQL queries?
197. How do you detect invalid packages affecting query performance?
198. How do you optimize multi-schema queries for complex reporting?
199. How do you monitor HADR logs for replay latency?
200. How do you perform historical performance analysis using snapshot tables?

🔷 Advanced Level (201–500)

Query Performance & Advanced Tuning (201–260)

201. How do you analyze expensive queries using db2top or db2pd?
202. How do you detect high CPU queries and optimize them?
203. How do you tune queries using EXPLAIN plan analysis?
204. How do you optimize queries with multiple table joins?
205. How do you rewrite correlated subqueries for better performance?
206. How do you detect queries causing sort heap overflows?
207. How do you implement temporary tables efficiently for complex queries?
208. How do you optimize GROUP BY and ORDER BY queries?
209. How do you optimize SELECT DISTINCT queries on large tables?
210. How do you use multi-dimensional indexes to speed up queries?
211. How do you detect full table scans and implement indexes?
212. How do you optimize queries on VARCHAR and CLOB columns?
213. How do you implement query hints for index usage?
214. How do you tune queries for partitioned tables?
215. How do you detect queries causing lock contention?
216. How do you implement index-only access for frequent queries?

Praveen Madupu
Mb: +91 98661 30093
Sr. Database Administrator

217. How do you optimize queries with IN or OR conditions?
218. How do you rewrite queries using EXISTS instead of IN for performance?
219. How do you monitor package cache usage for repeated queries?
220. How do you detect invalid packages and rebind them efficiently?
221. How do you tune queries with nested loops joins?
222. How do you optimize queries with multiple ORMs generating SQL?
223. How do you detect and resolve deadlocks in high-concurrency queries?
224. How do you tune DB2 sort heaps for large aggregation queries?
225. How do you optimize queries for read-only reporting workloads?
226. How do you analyze and optimize stored procedures?
227. How do you detect queries causing excessive log writes?
228. How do you optimize queries to reduce temporary tablespace usage?
229. How do you implement query optimization for star-schema OLAP queries?
230. How do you optimize queries with multiple CTEs (Common Table Expressions)?
231. How do you detect queries with inefficient data type conversions?
232. How do you optimize queries accessing large historical tables?
233. How do you tune index fetch strategies for high-performance OLTP queries?
234. How do you optimize queries causing heavy lock escalation?
235. How do you monitor query performance regression after DB2 upgrades?
236. How do you tune join order and join algorithms in complex queries?
237. How do you optimize queries on highly skewed datasets?
238. How do you implement parallelism for long-running queries?
239. How do you tune queries using materialized query tables (MQTs)?
240. How do you detect queries causing memory spills to disk?
241. How do you implement table partitioning to improve query performance?
242. How do you detect queries causing multiple buffer pool page faults?
243. How do you optimize OLAP queries for analytical dashboards?
244. How do you monitor query wait events to identify bottlenecks?
245. How do you tune queries for multi-core CPUs efficiently?
246. How do you optimize join queries using appropriate indexes?
247. How do you rewrite inefficient scalar functions in queries?
248. How do you implement automatic statistics collection for query optimization?
249. How do you detect queries causing high I/O latency?
250. How do you optimize queries on partitioned and multi-dimensional tables?
251. How do you analyze query performance trends over time?
252. How do you detect queries hitting slow storage?
253. How do you optimize queries using parallel index scans?
254. How do you tune queries with temporary LOB usage?
255. How do you detect queries causing log file contention?
256. How do you optimize queries accessing multiple schema objects?
257. How do you monitor query execution plans in production?
258. How do you optimize queries using approximate query techniques?
259. How do you implement workload balancing for read-heavy queries?

Praveen Madupu
Mb: +91 98661 30093
Sr. Database Administrator

260. How do you monitor expensive SQL statements continuously?

Indexing & Partitioning Advanced (261–300)

261. How do you create partitioned tables for performance and storage management?
262. How do you implement range, list, and hash partitioning in DB2?
263. How do you detect and rebuild fragmented indexes online?
264. How do you monitor index usage statistics?
265. How do you implement compressed indexes for storage optimization?
266. How do you create multi-column indexes for query optimization?
267. How do you detect redundant or unused indexes?
268. How do you implement index-only access plans?
269. How do you monitor index page splits and fragmentation?
270. How do you optimize indexing strategies for mixed OLTP/OLAP workloads?
271. How do you implement unique indexes to enforce data integrity?
272. How do you optimize foreign key lookups using indexes?
273. How do you monitor index usage during high-concurrency transactions?
274. How do you implement partition-level indexing?
275. How do you optimize indexes for large tables with frequent updates?
276. How do you rebuild indexes with minimal system impact?
277. How do you implement adaptive indexing for dynamic workloads?
278. How do you tune index fill factors for write-heavy workloads?
279. How do you detect queries not utilizing indexes properly?
280. How do you implement index strategies for time-series data?
281. How do you optimize composite indexes for multiple queries?
282. How do you detect indexes causing lock contention?
283. How do you optimize indexes for reporting queries on OLAP systems?
284. How do you implement automatic index monitoring scripts?
285. How do you tune indexes for partitioned table queries?
286. How do you implement concurrent index rebuilds?
287. How do you monitor index cache hit ratios?
288. How do you implement bitmap indexes for low-cardinality columns?
289. How do you detect slow index scans using performance monitors?
290. How do you optimize indexes for high-volume inserts?
291. How do you analyze index usage trends over time?
292. How do you implement index compression to save space?
293. How do you optimize indexes for star-schema queries?
294. How do you tune multi-level indexes for deep table hierarchies?
295. How do you implement partition pruning for query performance?
296. How do you monitor multi-column index efficiency?
297. How do you detect queries bypassing indexes due to data type mismatch?
298. How do you implement clustered indexes in DB2?
299. How do you optimize indexes in high-concurrency OLTP environments?
300. How do you balance index maintenance overhead versus query performance?

Praveen Madupu
Mb: +91 98661 30093
Sr. Database Administrator

HADR, High Availability & Replication (301–340)

301. How do you configure HADR primary and standby databases?
302. How do you monitor HADR synchronization status?
303. How do you implement synchronous vs asynchronous replication?
304. How do you detect replication lag and latency issues?
305. How do you perform manual role switch in HADR?
306. How do you add a new standby to an existing HADR setup?
307. How do you remove a failed standby safely?
308. How do you implement read-only standby for reporting?
309. How do you optimize HADR log shipping performance?
310. How do you automate HADR failover testing?
311. How do you monitor HADR logs for replay latency?
312. How do you detect HADR network connectivity issues?
313. How do you implement multi-site HADR deployments?
314. How do you handle split-brain scenarios in HADR?
315. How do you monitor transaction commit latency across nodes?
316. How do you configure HADR client redirection after failover?
317. How do you tune HADR for write-heavy workloads?
318. How do you implement backup strategies in HADR environments?
319. How do you monitor HADR standby database performance?
320. How do you detect synchronization failures during heavy transactions?
321. How do you implement multiple standby databases for redundancy?
322. How do you monitor HADR automatic reconnect behavior?
323. How do you optimize HADR for geographically distributed sites?
324. How do you detect transaction conflicts in multi-node replication?
325. How do you tune HADR parameters for minimal latency?
326. How do you perform point-in-time recovery in HADR setups?
327. How do you monitor HADR failover events?
328. How do you implement alerting for HADR role switches?
329. How do you monitor network bandwidth for HADR replication?
330. How do you optimize HADR log buffer size for high throughput?
331. How do you implement read scaling using standby databases?
332. How do you monitor HADR replication queue depth?
333. How do you detect standby database lagging during peak hours?
334. How do you perform emergency takeover in HADR setups?
335. How do you integrate HADR monitoring with Prometheus or Grafana?
336. How do you detect failed log shipping in HADR?
337. How do you optimize HADR synchronous writes on SSD storage?
338. How do you monitor HADR transaction rate per second?
339. How do you implement automatic failback strategies?
340. How do you optimize multi-standby HADR performance?

Monitoring, Security, Automation & Advanced Troubleshooting (341–400)

341. How do you monitor DB2 instance-level performance metrics continuously?
342. How do you implement automated alerts for CPU and memory spikes?
343. How do you detect sudden disk I/O bottlenecks?
344. How do you monitor buffer pool hit ratios in real-time?
345. How do you implement historical trend analysis of SQL performance?
346. How do you automate collection of snapshot and monitor data?
347. How do you detect queries consuming excessive temporary tablespace?
348. How do you monitor long-running transactions per application?
349. How do you detect lock contention between concurrent sessions?
350. How do you implement automated deadlock detection scripts?
351. How do you monitor package cache and rebind invalid packages automatically?
352. How do you automate index usage analysis and optimization?
353. How do you implement alerts for HADR standby lag?
354. How do you monitor replication conflicts in multi-site setups?
355. How do you implement automated database health checks?
356. How do you detect and prevent privilege escalation attacks?
357. How do you implement encryption at rest and in transit?
358. How do you audit user access and activities efficiently?
359. How do you detect unauthorized connections or failed login attempts?
360. How do you automate instance-level parameter adjustments?
361. How do you monitor high-volume log writes and detect log bottlenecks?
362. How do you automate backup verification and reporting?
363. How do you monitor and tune sort heap usage across multiple queries?
364. How do you detect temporary table contention in high-concurrency workloads?
365. How do you automate collection of query runtime statistics?
366. How do you implement automated alerting for tablespace full conditions?
367. How do you detect sudden growth in archive logs or active logs?
368. How do you monitor multi-node replication and connectivity health?
369. How do you automate failover testing in HADR environments?
370. How do you monitor application-specific query performance?
371. How do you implement continuous monitoring for high-priority workloads?
372. How do you detect and resolve excessive lock escalation automatically?
373. How do you monitor active utility operations and long-running maintenance tasks?
374. How do you implement automated reconciliation of data inconsistencies?
375. How do you monitor system resource usage per DB2 instance?
376. How do you implement automated log cleanup for long-running instances?
377. How do you detect high-latency queries on standby databases?
378. How do you monitor SQL statement frequency and detect abnormal spikes?
379. How do you automate notification for tablespace growth beyond thresholds?
380. How do you detect and alert on memory leaks in DB2 processes?
381. How do you monitor network latency impacting replication or queries?
382. How do you implement automated tuning of frequently executed queries?

Praveen Madupu
Mb: +91 98661 30093
Sr. Database Administrator

383. How do you monitor and optimize HADR synchronous commit performance?
384. How do you detect performance degradation after package binds?
385. How do you automate detection of invalid objects in the database?
386. How do you monitor temporary tablespace utilization per session?
387. How do you implement automated analysis for high-volume insert operations?
388. How do you monitor SQL queries causing excessive log generation?
389. How do you detect and prevent runaway queries automatically?
390. How do you implement automated index rebuild based on usage patterns?
391. How do you monitor multi-node transactions and detect blocking chains?
392. How do you implement automated alerts for replication failures?
393. How do you monitor and optimize memory usage across buffer pools?
394. How do you implement automated detection for table or index fragmentation?
395. How do you monitor the performance impact of new application deployments?
396. How do you detect and resolve contention on temporary LOB storage?
397. How do you automate analysis of slow-running reports or batch jobs?
398. How do you monitor and tune the database manager configuration parameters?
399. How do you implement automated snapshot collection for performance tuning?
400. How do you detect sudden degradation in HADR throughput?

Advanced Backup, Recovery & Disaster Recovery (401–450)

401. How do you implement incremental backup strategies for very large databases?
402. How do you monitor backup job performance in multi-instance setups?
403. How do you automate point-in-time recovery testing?
404. How do you restore individual tables or tablespaces from backups?
405. How do you implement log retention policies in production?
406. How do you monitor active log usage and prevent log full conditions?
407. How do you implement offsite backup rotation policies?
408. How do you detect corrupt backup sets before restoration?
409. How do you automate backup compression and encryption?
410. How do you implement hot backups with minimal downtime?
411. How do you perform restore operations in HADR primary and standby databases?
412. How do you automate backup verification for all databases?
413. How do you monitor incremental and delta backup performance?
414. How do you restore a database to an alternate location for testing?
415. How do you implement automated archive log shipping to standby systems?
416. How do you monitor long-running backup jobs and detect stalls?
417. How do you implement disaster recovery drills for large-scale databases?
418. How do you detect missing logs required for point-in-time recovery?
419. How do you implement continuous backup strategies in OLTP systems?
420. How do you optimize backup performance for very large tablespaces?
421. How do you implement mirrored backup strategies for redundancy?
422. How do you monitor backup storage consumption trends?
423. How do you automate restore and recovery scripts in DR environments?

424. How do you detect failures in automated backup jobs?
425. How do you implement cross-instance backup monitoring?
426. How do you monitor recovery progress in point-in-time scenarios?
427. How do you automate backup retention cleanup for old snapshots?
428. How do you optimize backup performance for high-concurrency environments?
429. How do you implement HADR-aware backup strategies?
430. How do you monitor backup job SLA adherence?
431. How do you automate reporting of backup failures?
432. How do you detect inconsistencies in restored databases?
433. How do you implement multi-site backup replication for disaster recovery?
434. How do you monitor archive log shipping performance?
435. How do you implement automated backup alerts for critical tablespaces?
436. How do you optimize backup for databases with heavy LOB usage?
437. How do you monitor backup job impact on production workloads?
438. How do you automate backup verification for compliance audits?
439. How do you restore large databases efficiently using parallelism?
440. How do you implement incremental backup verification scripts?
441. How do you monitor and optimize delta backup I/O throughput?
442. How do you implement automated cleanup of incomplete backups?
443. How do you detect slow restores and optimize restore time?
444. How do you implement backup redundancy across multiple storage devices?
445. How do you monitor backup jobs in HADR synchronous mode?
446. How do you automate validation of restored standby databases?
447. How do you implement multi-level backup strategies for critical systems?
448. How do you detect and alert for backup performance degradation?
449. How do you implement automated disaster recovery plan testing?
450. How do you optimize backup and recovery procedures for minimal downtime?

Extreme Performance Tuning & Troubleshooting (451–500)

451. How do you detect I/O bottlenecks affecting query performance?
452. How do you monitor buffer pool contention for high-concurrency workloads?
453. How do you tune instance-level memory for mixed OLTP and OLAP workloads?
454. How do you detect and resolve temporary tablespace bottlenecks?
455. How do you optimize queries causing high log writes?
456. How do you monitor and reduce lock escalation events?
457. How do you implement automatic workload balancing across multiple DB2 instances?
458. How do you detect performance regressions after DB2 upgrades or patches?
459. How do you optimize LOB storage for high-performance applications?
460. How do you implement dynamic tuning of buffer pools based on workload?
461. How do you monitor multi-node transaction performance in HADR clusters?
462. How do you detect SQL statements causing excessive CPU usage?
463. How do you implement query parallelism to speed up complex analytics?
464. How do you tune sort heaps for high-volume aggregation queries?

465. How do you monitor and optimize concurrent utility operations?
466. How do you detect and resolve deadlocks automatically?
467. How do you implement adaptive query optimization strategies?
468. How do you monitor and optimize network latency affecting HADR replication?
469. How do you detect and tune queries causing buffer pool churn?
470. How do you implement automated query rewriting for performance improvements?
471. How do you monitor long-running transactions impacting OLTP performance?
472. How do you detect and tune queries with suboptimal join methods?
473. How do you optimize queries on partitioned tables for minimal I/O?
474. How do you implement advanced indexing strategies for high-performance queries?
475. How do you monitor HADR lag during peak workloads?
476. How do you implement workload separation for OLTP and reporting queries?
477. How do you detect queries causing temp table spills to disk?
478. How do you implement proactive query tuning using historical stats?
479. How do you monitor and optimize log buffer usage for high throughput?
480. How do you detect tablespace hot spots and redistribute data?
481. How do you implement multi-buffer pool strategy for mixed workloads?
482. How do you monitor package cache and prevent invalidations?
483. How do you detect and resolve index contention issues?
484. How do you implement automatic tuning of frequently used queries?
485. How do you monitor and optimize join performance for star-schema databases?
486. How do you detect performance anomalies using historical baselines?
487. How do you implement query plan stability strategies for production workloads?
488. How do you monitor resource consumption per application user?
489. How do you detect and optimize SQL causing excessive sort or group operations?
490. How do you implement automated alerts for query SLA violations?
491. How do you monitor HADR transaction replay for consistency?
492. How do you detect memory leaks in DB2 agents or utilities?
493. How do you optimize OLAP queries on partitioned and compressed tables?
494. How do you implement advanced backup strategies for mission-critical databases?
495. How do you detect storage bottlenecks impacting DB2 performance?
496. How do you monitor temp LOB storage utilization?
497. How do you implement automated remediation for query performance regressions?
498. How do you monitor and optimize SQL parallelism settings?
499. How do you detect excessive lock waits and implement automatic mitigation?
500. How do you implement full end-to-end monitoring, tuning, and recovery strategies for DB2 LUW in production environments?