Praveen Madupu
Mb: +91 98661 30093
Sr. Database Administrator

Real-time scenario-based interview questions for **Amazon DynamoDB Administration, Performance Tuning & Optimization**

🟢 **Amazon DynamoDB Administration, Performance Tuning & Optimization – 500 Real-Time Scenario Questions**

🔷 **Basic Level (0–2 years) – 1–100**

DynamoDB Setup & Configuration

1. How do you create a DynamoDB table using the AWS Management Console?
2. How do you choose between on-demand and provisioned capacity modes?
3. How do you define a primary key (partition key) for a table?
4. How do you define a composite primary key (partition key + sort key)?
5. How do you enable DynamoDB Streams on a table?
6. How do you configure TTL (Time to Live) for expiring items automatically?
7. How do you enable encryption at rest for a DynamoDB table?
8. How do you configure encryption in transit for DynamoDB operations?
9. How do you define Global Secondary Indexes (GSI) for querying efficiently?
10. How do you define Local Secondary Indexes (LSI) for range queries?

Table & Capacity Management

11. How do you scale a DynamoDB table with on-demand capacity?
12. How do you adjust read/write capacity units for provisioned tables?
13. How do you add a new Global Secondary Index to an existing table?
14. How do you remove a Global Secondary Index safely?
15. How do you monitor consumed read/write capacity in real-time?
16. How do you monitor throttled requests for a table?
17. How do you enable auto-scaling for read/write capacity?
18. How do you detect hot partitions in a high-traffic table?
19. How do you optimize partition key design to avoid throttling?
20. How do you rename a table or index in DynamoDB?

Security & User Management

21. How do you configure IAM policies for DynamoDB access?
22. How do you restrict access to specific tables or indexes?
23. How do you enforce fine-grained access control at the item level?
24. How do you monitor unauthorized access attempts?
25. How do you enable encryption using a customer-managed KMS key?
26. How do you rotate encryption keys for DynamoDB tables?
27. How do you enforce HTTPS for all DynamoDB API calls?
28. How do you integrate DynamoDB access with AWS Cognito?
29. How do you audit all DynamoDB API calls using CloudTrail?
30. How do you monitor IAM roles used for DynamoDB operations?

Backup & Recovery Basics

31. How do you perform on-demand backups of a DynamoDB table?

32. How do you restore a DynamoDB table from a backup?
33. How do you enable continuous backups with point-in-time recovery (PITR)?
34. How do you monitor backup completion status?
35. How do you schedule automated backups for multiple tables?
36. How do you perform cross-region table backups?
37. How do you restore a single table from PITR?
38. How do you verify the integrity of backups?
39. How do you monitor backup storage costs?
40. How do you optimize backup frequency without impacting performance?

Monitoring & Logging Basics

41. How do you enable CloudWatch metrics for DynamoDB?
42. How do you monitor read/write capacity usage per table?
43. How do you monitor throttled requests and identify causes?
44. How do you track consumed vs. provisioned capacity in real-time?
45. How do you enable detailed CloudWatch metrics for hot partitions?
46. How do you monitor DynamoDB Streams processing?
47. How do you detect anomalies in query latency?
48. How do you track item size and payload trends over time?
49. How do you set up CloudWatch alarms for capacity thresholds?
50. How do you enable logging for Lambda triggers on DynamoDB Streams?

Table Schema & Query Basics

51. How do you design a table schema for a high-throughput application?
52. How do you choose between GSI and LSI for query flexibility?
53. How do you query items efficiently using partition and sort keys?
54. How do you perform conditional updates to prevent overwrites?
55. How do you handle large items exceeding 400 KB in DynamoDB?
56. How do you implement composite keys for time-series data?
57. How do you perform batch writes efficiently?
58. How do you handle duplicate items during batch inserts?
59. How do you design a schema for multi-tenant applications?
60. How do you monitor and optimize query efficiency for large tables?

Query Performance Basics

61. How do you detect and reduce throttled queries?
62. How do you optimize partition key selection for read-heavy workloads?
63. How do you tune GSI and LSI for faster queries?
64. How do you avoid hot keys causing uneven load distribution?
65. How do you optimize scan operations for large datasets?
66. How do you minimize latency in high-concurrency queries?
67. How do you use Query vs. Scan effectively for different scenarios?
68. How do you detect and optimize slow query patterns?

Praveen Madupu
Mb: +91 98661 30093
Sr. Database Administrator

69. How do you handle queries requiring multiple indexes?
70. How do you optimize queries for time-range filters in sort keys?

Storage & Item Size Management

71. How do you monitor item size distribution in a table?
72. How do you optimize storage for large binary items?
73. How do you implement item versioning for conflict resolution?
74. How do you detect tables approaching storage limits?
75. How do you monitor storage costs for high-volume tables?
76. How do you split large payloads across multiple items efficiently?
77. How do you implement compression to reduce storage usage?
78. How do you monitor data growth trends for capacity planning?
79. How do you detect unnecessary item duplication?
80. How do you implement TTL for automatic cleanup of old items?

Security Basics

81. How do you enforce IAM-based access for multiple teams?
82. How do you detect unauthorized API calls in DynamoDB?
83. How do you integrate DynamoDB access with AWS Cognito roles?
84. How do you implement attribute-level access control?
85. How do you monitor failed write attempts due to access policies?
86. How do you audit access patterns for compliance?
87. How do you rotate encryption keys periodically?
88. How do you enforce server-side encryption for all tables?
89. How do you implement fine-grained access control for developers?
90. How do you monitor encryption key usage and access logs?

High Availability & Replication

91. How do you enable DynamoDB global tables for multi-region replication?
92. How do you handle conflicts in global tables?
93. How do you monitor replication lag across regions?
94. How do you detect failed writes in multi-region deployments?
95. How do you implement cross-region disaster recovery for DynamoDB?
96. How do you optimize write throughput for global tables?
97. How do you monitor replication metrics using CloudWatch?
98. How do you detect anomalies in replication consistency?
99. How do you implement failover for multi-region tables?
100. How do you test disaster recovery scenarios safely?

Query Optimization & Indexing (101–180)

101. How do you choose the right partition key to minimize hot partitions?
102. How do you design a sort key for time-series data?
103. How do you decide between using GSI and LSI for queries?

104. How do you implement composite keys for multi-attribute querying?
105. How do you monitor and optimize GSI performance under load?
106. How do you detect and fix slow queries due to missing indexes?
107. How do you optimize scan operations to avoid full table scans?
108. How do you use query filters efficiently without increasing read capacity?
109. How do you perform batch queries for large datasets efficiently?
110. How do you prevent GSI throttling in high-write tables?
111. How do you detect queries causing excessive read capacity consumption?
112. How do you implement secondary indexes for multiple access patterns?
113. How do you monitor index usage trends over time?
114. How do you optimize queries for sparse attributes in GSIs?
115. How do you handle large queries that return paginated results?
116. How do you tune Query API requests for low latency?
117. How do you design a table schema for multiple query patterns?
118. How do you detect inefficient queries using CloudWatch metrics?
119. How do you optimize query performance for large text attributes?
120. How do you handle queries requiring multiple index lookups?
121. How do you design a schema for range queries on numerical data?
122. How do you optimize queries for frequent reads of small items?
123. How do you detect queries causing hot key throttling?
124. How do you implement efficient pagination for large datasets?
125. How do you monitor read/write request distribution across partitions?
126. How do you optimize queries for date-range filters using sort keys?
127. How do you design a schema for multi-tenant applications?
128. How do you monitor slow queries on GSIs?
129. How do you detect queries that bypass indexes?
130. How do you implement conditional writes for data integrity?
131. How do you optimize batch get operations for large key sets?
132. How do you design GSIs for analytical queries without impacting performance?
133. How do you implement secondary indexes for frequently updated attributes?
134. How do you monitor query latency using CloudWatch custom metrics?
135. How do you detect and fix queries with high network latency?
136. How do you optimize queries for sorted results without scanning all items?
137. How do you detect inefficient use of filters instead of keys?
138. How do you implement a caching layer to improve query latency?
139. How do you optimize queries for items with multiple attributes?
140. How do you monitor query efficiency using AWS X-Ray?
141. How do you detect queries causing disproportionate read consumption?
142. How do you optimize queries for secondary indexes with sparse attributes?
143. How do you design table schemas for multi-region queries?
144. How do you implement query retry strategies for throttled requests?
145. How do you detect and resolve hot partitions caused by skewed keys?
146. How do you optimize batch write operations for multiple tables?

147. How do you monitor query patterns to optimize index usage?
148. How do you detect queries causing scan-based bottlenecks?
149. How do you design tables for frequently accessed attributes?
150. How do you optimize queries for global tables without increasing latency?
151. How do you detect queries causing excessive conditional checks?
152. How do you monitor the impact of large queries on provisioned capacity?
153. How do you optimize GSIs for write-heavy workloads?
154. How do you detect slow queries caused by inefficient sort keys?
155. How do you design composite keys for hierarchical data models?
156. How do you monitor query latency trends for performance tuning?
157. How do you optimize queries for multi-attribute filtering?
158. How do you implement query throttling mitigation strategies?
159. How do you detect queries causing network congestion between AWS regions?
160. How do you optimize queries for time-based aggregation?
161. How do you monitor query performance using AWS CloudTrail logs?
162. How do you detect inefficient queries affecting other tables?
163. How do you optimize queries for sparse datasets?
164. How do you implement efficient query pagination for large datasets?
165. How do you detect queries causing high CPU utilization on DynamoDB Streams?
166. How do you optimize queries for attributes with high cardinality?
167. How do you monitor query execution patterns for analytics?
168. How do you detect queries causing duplicate data reads?
169. How do you optimize queries for multi-tenant partitioning?
170. How do you design schema for sorted key queries across multiple attributes?
171. How do you monitor query response times across different endpoints?
172. How do you detect inefficient use of scan filters on large tables?
173. How do you optimize queries for frequently updated attributes?
174. How do you implement efficient conditional writes for multi-user updates?
175. How do you detect queries causing throughput spikes on GSIs?
176. How do you optimize queries for multiple partition key patterns?
177. How do you monitor query latency during peak load?
178. How do you detect and fix queries causing partition hot spots?
179. How do you optimize queries for time-series analytics using sort keys?
180. How do you implement schema changes without downtime?

Global Tables & Multi-Region Replication (181–250)

181. How do you enable DynamoDB global tables for multi-region replication?
182. How do you detect replication conflicts in global tables?
183. How do you monitor replication latency across AWS regions?
184. How do you troubleshoot inconsistent data in global tables?
185. How do you implement conflict resolution strategies?
186. How do you optimize write throughput for global tables?
187. How do you monitor global table replication metrics in CloudWatch?

188. How do you detect replication lag during peak workloads?
189. How do you perform failover in a multi-region setup?
190. How do you test disaster recovery in a global table deployment?
191. How do you handle cross-region latency in high-traffic applications?
192. How do you implement multi-region writes safely?
193. How do you monitor multi-region table consistency?
194. How do you detect conflicts between region updates?
195. How do you resolve write conflicts using last-writer-wins strategy?
196. How do you optimize global table reads to reduce latency?
197. How do you monitor global table replication status continuously?
198. How do you detect failed replication writes automatically?
199. How do you optimize GSIs in global tables for minimal replication lag?
200. How do you implement schema changes across all regions?
201. How do you detect region-specific throttling issues?
202. How do you monitor multi-region write capacity utilization?
203. How do you optimize network traffic between global table regions?
204. How do you perform cross-region PITR restores?
205. How do you detect replication conflicts in multi-master setups?
206. How do you monitor read/write distribution across global table regions?
207. How do you optimize conflict resolution rules for business-critical data?
208. How do you monitor hot partitions in global tables?
209. How do you detect and resolve delayed replication in a specific region?
210. How do you implement automated alerts for replication failures?
211. How do you optimize query patterns to avoid cross-region latency?
212. How do you test failover scenarios for global tables without downtime?
213. How do you monitor multi-region throughput usage?
214. How do you detect regions experiencing write spikes?
215. How do you implement automated failback after disaster recovery?
216. How do you optimize global table partition keys for balanced writes?
217. How do you monitor replication errors caused by schema changes?
218. How do you detect high replication latency affecting analytics queries?
219. How do you optimize global tables for read-heavy workloads?
220. How do you monitor consistency checks across regions?
221. How do you detect and resolve write conflicts on GSIs in global tables?
222. How do you implement automated alerts for cross-region replication issues?
223. How do you optimize query patterns for multi-region applications?
224. How do you monitor endpoint health for global table regions?
225. How do you detect region-specific performance degradation?
226. How do you optimize global table write capacity during peak hours?
227. How do you implement cross-region table migrations safely?
228. How do you monitor multi-region table backups?
229. How do you detect replication delays caused by large batch writes?
230. How do you optimize global table reads to minimize latency?

231. How do you implement schema evolution in multi-region tables?
232. How do you detect hot keys affecting replication consistency?
233. How do you optimize conflict resolution for concurrent updates?
234. How do you monitor replication metrics using AWS CloudWatch dashboards?
235. How do you detect inconsistencies caused by eventual consistency?
236. How do you optimize global tables for write-heavy workloads?
237. How do you implement alerts for partition-level replication lag?
238. How do you monitor data integrity across all regions?
239. How do you detect replication throttling during high traffic periods?
240. How do you optimize queries to use local endpoints in global tables?
241. How do you implement cross-region backup strategies?
242. How do you monitor multi-region GSIs for replication performance?
243. How do you detect and fix failed cross-region updates?
244. How do you optimize query latency for global table analytics?
245. How do you monitor replication throughput trends?
246. How do you detect region-specific throttled requests in global tables?
247. How do you implement automated testing for replication consistency?
248. How do you optimize partition key selection for global tables?
249. How do you monitor replication conflicts in near real-time?
250. How do you implement multi-region failover for mission-critical tables?

Backup, PITR & Recovery (251–320)

251. How do you enable point-in-time recovery (PITR) on a DynamoDB table?
252. How do you restore a table to a specific timestamp using PITR?
253. How do you monitor the progress of a PITR restore operation?
254. How do you schedule periodic on-demand backups without impacting production?
255. How do you restore a table in a different AWS region?
256. How do you validate data integrity after a restore?
257. How do you restore a global table in case of cross-region failures?
258. How do you perform a PITR restore for large tables efficiently?
259. How do you monitor backup storage costs across multiple tables?
260. How do you implement automated backup verification scripts?
261. How do you restore only specific items from a backup?
262. How do you monitor failed or incomplete backup jobs?
263. How do you optimize backup throughput for large datasets?
264. How do you implement backup retention policies for compliance?
265. How do you restore tables to a new table name without overwriting production?
266. How do you perform cross-account backup and restore?
267. How do you test disaster recovery procedures safely?
268. How do you monitor PITR-enabled tables for unexpected changes?
269. How do you handle failed restores due to insufficient permissions?
270. How do you optimize backup strategies for high-write tables?
271. How do you restore large tables with minimal downtime?

272. How do you monitor incremental backups for errors?
273. How do you detect failed snapshot jobs and automate retries?
274. How do you restore a table to an earlier schema version?
275. How do you implement automated backup rotation?
276. How do you monitor backup throughput and latency?
277. How do you restore a table to a region with different read/write capacity?
278. How do you optimize PITR restores for multi-region global tables?
279. How do you detect and mitigate backup-induced throttling?
280. How do you monitor cross-region backup replication health?
281. How do you restore a table without impacting ongoing writes?
282. How do you automate verification of restored data integrity?
283. How do you monitor backup size trends for capacity planning?
284. How do you implement retention limits for long-term backups?
285. How do you detect partial backup failures in large tables?
286. How do you optimize restore performance using parallel operations?
287. How do you restore only selected attributes from a backup?
288. How do you monitor backup and restore operations for SLA compliance?
289. How do you handle backup restore conflicts in global tables?
290. How do you implement automated alerts for failed PITR restores?
291. How do you detect and recover from accidental deletes using PITR?
292. How do you monitor backup job duration for performance tuning?
293. How do you restore tables from multiple incremental backups efficiently?
294. How do you detect backup bottlenecks in high-write workloads?
295. How do you optimize cross-region restores for minimal latency?
296. How do you monitor restore operations using CloudWatch metrics?
297. How do you restore tables in a different AWS account?
298. How do you perform automated validation after restores?
299. How do you monitor long-term backup compliance?
300. How do you implement automated PITR restore testing?
301. How do you detect inconsistent data after restore?
302. How do you optimize restore operations for tables with GSIs?
303. How do you monitor restore throttling to avoid capacity issues?
304. How do you restore a table with encrypted items safely?
305. How do you optimize backup frequency for cost efficiency?
306. How do you monitor incremental backup storage growth?
307. How do you restore a table to a new AWS region endpoint?
308. How do you automate notifications for backup job completions?
309. How do you handle restore failures caused by schema mismatches?
310. How do you optimize backup throughput without impacting live traffic?
311. How do you restore tables with large payloads efficiently?
312. How do you monitor PITR for multi-region global tables?
313. How do you detect anomalies in backup job durations?
314. How do you implement automated restore verification scripts?

315. How do you restore tables under high read/write load?
316. How do you optimize backup storage by compressing items?
317. How do you monitor restore latency for SLA compliance?
318. How do you restore tables with heavy GSI usage?
319. How do you implement automated PITR restore for multiple tables?
320. How do you detect and resolve backup-induced throttling in large tables?

Performance Tuning & Monitoring (321–400)

321. How do you monitor read/write capacity usage in real-time?
322. How do you detect hot partitions affecting performance?
323. How do you optimize provisioned throughput for read-heavy workloads?
324. How do you optimize write-heavy tables using partition keys?
325. How do you monitor throttled requests using CloudWatch metrics?
326. How do you tune GSI write capacity to reduce throttling?
327. How do you optimize Query vs. Scan usage?
328. How do you monitor item size and its effect on performance?
329. How do you detect latency spikes in high-concurrency queries?
330. How do you optimize batch write operations to reduce throttling?
331. How do you monitor DynamoDB Streams for performance bottlenecks?
332. How do you tune global table replication performance?
333. How do you optimize read-heavy queries on GSIs?
334. How do you detect queries causing uneven partition usage?
335. How do you optimize partition key design for write-intensive workloads?
336. How do you monitor DynamoDB metrics using CloudWatch dashboards?
337. How do you detect and optimize queries consuming excess RCUs?
338. How do you tune LSI and GSI to avoid hot partitions?
339. How do you implement caching to reduce DynamoDB read latency?
340. How do you monitor throughput spikes and adjust capacity dynamically?
341. How do you detect throttled batch writes in large tables?
342. How do you optimize query patterns for time-series data?
343. How do you monitor latency trends over peak and off-peak periods?
344. How do you tune provisioned throughput to minimize costs?
345. How do you detect high-latency queries using CloudWatch Logs?
346. How do you optimize conditional writes for minimal conflicts?
347. How do you monitor read/write capacity utilization per application?
348. How do you detect inefficient scan operations?
349. How do you optimize table schemas to minimize throttling?
350. How do you monitor multi-region replication throughput for global tables?
351. How do you detect inefficient attribute selection in queries?
352. How do you optimize queries to minimize consumed RCUs/WCUs?
353. How do you tune auto-scaling settings to handle workload spikes?
354. How do you monitor GSI read/write activity under peak load?
355. How do you detect queries causing hot partitions in real-time?

356. How do you optimize queries for large datasets using pagination?
357. How do you monitor and optimize DynamoDB Streams latency?
358. How do you detect write-heavy partitions causing throttling?
359. How do you implement adaptive capacity for uneven workloads?
360. How do you optimize query patterns for frequently updated items?
361. How do you monitor application-level performance using X-Ray with DynamoDB?
362. How do you detect inefficient scan operations affecting throughput?
363. How do you optimize batch get operations for multiple keys?
364. How do you monitor hot key distribution trends over time?
365. How do you detect cross-region latency affecting global tables?
366. How do you optimize queries for large payload attributes?
367. How do you tune GSI/WCU to improve write performance?
368. How do you monitor partition-level performance metrics?
369. How do you detect spikes in conditional check failures?
370. How do you optimize write-heavy workloads for minimal throttling?
371. How do you monitor throttling trends to prevent SLA violations?
372. How do you detect inefficient usage of scan filters?
373. How do you optimize queries for frequent time-range queries?
374. How do you monitor cross-region replication lag for analytics?
375. How do you detect queries causing excessive network calls?
376. How do you optimize queries for multi-tenant partitioning?
377. How do you monitor read/write ratios for balanced workloads?
378. How do you detect queries causing high latency on GSIs?
379. How do you optimize table schemas for multi-access patterns?
380. How do you monitor query efficiency over time for trending analysis?
381. How do you detect inefficient use of attributes in queries?
382. How do you optimize queries for heavy analytical workloads?
383. How do you monitor DynamoDB Streams consumer latency?
384. How do you detect and fix partitions causing throughput imbalance?
385. How do you optimize global tables for cross-region analytics?
386. How do you monitor capacity consumption for multiple applications?
387. How do you detect queries bypassing indexes?
388. How do you optimize batch writes for high-concurrency scenarios?
389. How do you monitor endpoint performance for low-latency access?
390. How do you detect queries causing redundant data scans?
391. How do you optimize queries for high-cardinality attributes?
392. How do you monitor DynamoDB performance metrics in near real-time?
393. How do you detect queries causing CPU spikes in application logic?
394. How do you optimize query patterns for heavy read/write workloads?
395. How do you monitor query latency for SLA compliance?
396. How do you detect inefficient query pagination?
397. How do you optimize queries for frequently updated datasets?
398. How do you monitor consumed vs. provisioned capacity in global tables?

399. How do you detect queries causing write conflicts in multi-region tables?
400. How do you tune tables and indexes for extreme performance workloads?

Security, Hot Partitions & Extreme Scenarios (401–500)

401. How do you enforce fine-grained IAM policies for multiple teams?
402. How do you monitor unauthorized DynamoDB API access?
403. How do you implement attribute-level security for sensitive data?
404. How do you rotate KMS keys for encrypted tables?
405. How do you detect and mitigate hot partitions under extreme load?
406. How do you monitor AWS WCU/RCU usage for security-sensitive tables?
407. How do you optimize queries to avoid hot key throttling?
408. How do you detect sudden spikes in write-heavy workloads?
409. How do you implement partition key design for balanced global write load?
410. How do you monitor and alert on capacity saturation?
411. How do you detect replication conflicts affecting sensitive tables?
412. How do you optimize GSI usage for minimal latency under extreme load?
413. How do you implement automated scaling for provisioned capacity?
414. How do you detect and resolve multi-region consistency issues?
415. How do you monitor slow queries affecting SLA compliance?
416. How do you detect queries causing network congestion?
417. How do you implement caching strategies to reduce latency?
418. How do you monitor table performance under multi-tenant scenarios?
419. How do you detect and mitigate throttled batch operations?
420. How do you optimize queries for items exceeding 400 KB?
421. How do you monitor partition-level write spikes in real-time?
422. How do you detect inefficient query patterns in production workloads?
423. How do you optimize queries for high-throughput streaming data?
424. How do you implement cross-region failover for global tables?
425. How do you monitor latency and errors for DynamoDB Streams consumers?
426. How do you detect queries causing excessive RCUs/WCUs consumption?
427. How do you optimize GSI queries for write-heavy workloads?
428. How do you monitor and alert on hot partition detection?
429. How do you detect and prevent performance degradation during peak hours?
430. How do you optimize partition key selection for multi-attribute access?
431. How do you monitor slow queries impacting analytics workloads?
432. How do you detect write contention in multi-region global tables?
433. How do you implement automated query optimization alerts?
434. How do you monitor partition key distribution in real-time?
435. How do you detect queries causing throttling during batch operations?
436. How do you optimize queries for time-series analytics?
437. How do you monitor replication consistency for SLA compliance?
438. How do you detect inefficient scans affecting high-volume tables?
439. How do you implement schema changes without downtime in production?

440. How do you monitor cross-region latency for global applications?
441. How do you detect queries causing partition hot spots in global tables?
442. How do you optimize conditional updates for multi-user workloads?
443. How do you monitor and tune global table replication throughput?
444. How do you detect inefficient GSI usage causing latency spikes?
445. How do you implement automated alerts for failed queries?
446. How do you monitor DynamoDB Streams for delayed processing?
447. How do you detect queries causing disproportionate RCU/WCU usage?
448. How do you optimize queries for extreme write-heavy workloads?
449. How do you monitor table throughput for SLA violations?
450. How do you detect queries causing network congestion across regions?
451. How do you implement caching for extreme performance scenarios?
452. How do you monitor latency and throttling trends continuously?
453. How do you detect hot keys causing capacity spikes?
454. How do you optimize queries for analytical dashboards?
455. How do you monitor replication lag in near real-time for critical tables?
456. How do you detect and mitigate slow query performance during peak traffic?
457. How do you optimize queries for multi-attribute filtering under heavy load?
458. How do you monitor endpoint health and throttling metrics continuously?
459. How do you detect and fix queries causing global table inconsistencies?
460. How do you optimize batch write and read operations under extreme load?
461. How do you monitor and tune GSIs for multi-region applications?
462. How do you detect queries bypassing indexes causing performance degradation?
463. How do you implement schema evolution in extreme-scale tables?
464. How do you monitor partition-level performance metrics for SLA compliance?
465. How do you detect inefficient query patterns affecting replication?
466. How do you optimize queries for multi-tenant partitioning at scale?
467. How do you monitor throughput consumption for global tables?
468. How do you detect conditional write failures affecting high-volume tables?
469. How do you optimize queries for large payload items without throttling?
470. How do you monitor and detect cross-region replication conflicts in real-time?
471. How do you implement adaptive capacity for uneven workloads?
472. How do you detect hot partitions causing SLA breaches?
473. How do you optimize queries for frequent scan operations?
474. How do you monitor DynamoDB Streams consumer lag for critical processing?
475. How do you detect queries causing high network latency between regions?
476. How do you optimize GSIs for write-heavy and read-heavy mixed workloads?
477. How do you monitor RCU/WCU consumption trends for capacity planning?
478. How do you detect queries causing disproportionate impact on other tenants?
479. How do you implement automated query throttling mitigation strategies?
480. How do you monitor real-time latency for critical operations?
481. How do you detect and fix queries causing uneven global table replication?
482. How do you optimize batch operations for extreme concurrency?

483. How do you monitor query execution efficiency over large datasets?
484. How do you detect high-latency queries caused by inefficient GSI usage?
485. How do you optimize queries for multi-attribute and composite key patterns?
486. How do you monitor partition hot spots dynamically under varying load?
487. How do you detect queries causing excessive cross-region data transfer?
488. How do you optimize conditional updates and deletes for extreme load?
489. How do you monitor query efficiency and RCU/WCU usage for SLA enforcement?
490. How do you detect replication lag affecting analytics dashboards?
491. How do you optimize global tables for high availability and low latency?
492. How do you monitor throughput distribution across partitions?
493. How do you detect queries causing replication conflicts in GSIs?
494. How do you optimize queries to minimize read/write consumption cost?
495. How do you monitor global table replication metrics for performance tuning?
496. How do you detect inefficient query and scan patterns in production workloads?
497. How do you optimize DynamoDB Streams for extreme throughput workloads?
498. How do you monitor partition utilization to prevent SLA breaches?
499. How do you detect and fix conditional write failures affecting global tables?
500. How do you implement extreme performance tuning for global tables while maintaining HA, minimal latency, and cost efficiency?